

CLOUD / DISTRIBUTED SYSTEMS: HANDLING CONCURRENCY / SYNCHRONIZATION WITH DB LOCKING

Andrew Rembrandt,
andrew@3ap.ch

TechHub Swiss
24 February, 2022

3AP Lightning talk / 'Brain snack'
14 April, 2022

WHEN & METHODS OF SYNCHRONISATION / LOCKING

- Why/when
 - Multiple instances of a service (cloud / HA / ...)
 - Exclusive locking scenario
 - Unable to implement ordered events / requests
 - Cloud pubsub Ordering Keys
 - Kafka Partitioning

MULTI-PROCESS LOCKING

- Datastore
 - DB (Postgres, Oracle ...)
 - NoSQL (Redis, Zookeeper, Hazelcast, MongoDB, ...)

DB LOCKING

- Pessimistic lock
 - Row / table level => contention / performance impact

```
@Lock(LockModeType.PESSIMISTIC_WRITE)
@QueryHints({@QueryHint(name = "javax.persistence.lock.timeout",
value = "1000")})
@Query("SELECT p FROM person p WHERE p.id = ?1")
Optional<PersonEntity> findByIdPessimisticLocked(Long id);
```

- Optimistic locking
 - Atomic compare and set with a version or timestamp

```
@Lock(LockModeType.OPTIMISTIC_FORCE_INCREMENT)
@Query("FROM shared_lock WHERE lockType = :lockType AND sourceId =
:sourceId")
Optional<SharedLockEntity> findByLockTypeAndSourceIdLocked(LockType
lockType, Long sourceId);
```

OPTIMISTIC LOCKING (CONT)

- Typical sql statement:

```
update
  item
set
  version=1,
  amount=10
where
  id='abcd1234'
and
  version=0
```

- Existing entity locking (version / lock timestamp column)
- Dedicated lock table
 - Decoupling provides more flexibility => 'multi-process mutex'

REACTIVE EXAMPLE

```
public Mono<FolioEntity> createOrUpdateFolioTransaction(Event event)
{
    return doCreateOrUpdateFolioTransaction(event)
        .retryWhen(
            Retry.backoff(5, Duration.ofMillis(300))
        )
        .filter(OptimisticLockingFailureException.class::isInstance)
        .doBeforeRetry(
            s ->
                log.warn("Optimistic locking failure while
createOrUpdateFolioTransaction for ...", ...)
        )
}

private Mono<FolioEntity> doCreateOrUpdateFolioTransaction(Event
event) {
    return Mono.defer(() -> {
        FolioEntity folioEntity = null;
        Optional<FolioEntity> folio =

folioRepository.findByFolioIdLocked(event.getObjectId());

        if (folio.isEmpty()) {
            folioEntity = mapper.createFolioEntity(event);
        } else {
            folioEntity = mapper.updateFolioEntity(event,
folio.get());
        }
    });
}
```

```
        }  
        return folioRepository.save(folioEntity);  
    })  
    .flatMap(folioRepositoryService::addTransientFields);  
}
```

GENERIC DB SHARED LOCK TABLE

```
public <T> Mono<T> acquireLock(Long sourceId, String reportableId,
    LockType lockType, Supplier<Mono<T>> callbackOnLockAcquisition) {
    return Mono.defer(
        () ->
            sharedLockRepository
                .findByLockTypeAndSourceIdLocked(lockType,
sourceId)
                .or((() -> Optional.of(
sharedLockRepository.save(SharedLockEntity.builder()
    .lockType(lockType).sourceId(sourceId).failureCount(0)
                                .expiryCount(0).build()))))
                .flatMap(
                    lock -> {
                        if (lock.getAcquiredLockTime() != null)
return Optional.empty();
                        lock.setAcquiredLockTime(Instant.now());
                        val updatedLock =
sharedLockRepository.save(lock);
                        return Optional.of(updatedLock);
                    })
                .onErrorResume( // OpLock exeception for existing rows,
integrity violation for conflicting row inserts
                    t -> t instanceof ObjectOptimisticLockingFailureException
```



```
|| t instanceof DataIntegrityViolationException,  
    e -> Mono.empty()  
    })  
    .flatMap(lock -> callbackOnLockAcquisition.get()) // We have  
acquired the lock, call the callback  
    .flatMap(  
        resultItem ->  
            tenantContext.fromSupplier(  
                () -> {  
                    log.info("Remove lock for {}, sourceId: {}",  
reportableId, sourceId);  
                    sharedLockRepository.unsetLock(lockType,  
sourceId);  
                    return resultItem;  
                })  
            )));  
}
```

5 MINS != EXPERT :)

- To lock or not to lock
 - Complexity trade-off
 - DB locks are simple until it's use grows
 - Idempotent updates
 - Event normalisation
- Audit trail / logging

QUESTIONS

- **<https://github.com/andrewrembrandt/distributed-locks-talk>**
- **<https://blog.mimacom.com/tag/concurrency/>**