

# **CLOUD / DISTRIBUTED SYSTEMS: HANDLING CONCURRENCY / SYNCHRONIZATION WITH DB LOCKING**

Andrew Rembrandt,  
**andrew@3ap.ch**

TechHub Swiss  
24 February, 2022

3AP Lightning talk / 'Brain snack'  
14 April, 2022

## **WHEN & METHODS OF SYNCHRONISATION / LOCKING**

- Why/when
  - Multiple instances of a service (cloud / HA / ...)
  - Exclusive locking scenario
  - Unable to implement ordered events / requests
    - Cloud pubsub Ordering Keys
    - Kafka Partitioning

## **MULTI-PROCESS LOCKING**

- Datastore
  - DB (Postgres, Oracle ...)
  - NoSQL (Redis, Zookeeper, Hazelcast, MongoDB, ...)

## DB LOCKING

- Pessimistic lock
  - Row / table level => contention / performance impact

```
@Lock(LockModeType.PESSIMISTIC_WRITE)
@QueryHints({@QueryHint(name = "javax.persistence.lock.timeout",
value = "1000")})
@Query("SELECT p FROM person p WHERE p.id = ?1")
Optional<PersonEntity> findByIdPessimisticLocked(Long id);
```

- Optimistic locking
  - Atomic compare and set with a version or timestamp

```
@Lock(LockModeType.OPTIMISTIC_FORCE_INCREMENT)
@Query("FROM shared_lock WHERE lockType = :lockType AND sourceId =
:sourceId")
Optional<SharedLockEntity> findByLockTypeAndSourceIdLocked(LockType
lockType, Long sourceId);
```

## OPTIMISTIC LOCKING (CONT)

- Typical sql statement:

```
update
  item
set
  version=1,
  amount=10
where
  id='abcd1234'
and
  version=0
```

- Existing entity locking (version / lock timestamp column)
- Dedicated lock table
  - Decoupling provides more flexibility => 'multi-process mutex'

## REACTIVE EXAMPLE

```
public Mono<FolioEntity> createOrUpdateFolioTransaction(Event event)
{
    return doCreateOrUpdateFolioTransaction(event)
        .retryWhen(
            Retry.backoff(5, Duration.ofMillis(300))

        .filter(OptimisticLockingFailureException.class::isInstance)
            .doBeforeRetry(
                s ->
                    log.warn("Optimistic locking failure while
createOrUpdateFolioTransaction for ...", ...))
            }

private Mono<FolioEntity> doCreateOrUpdateFolioTransaction(Event
event) {
    return Mono.defer(() -> {
        FolioEntity folioEntity = null;
        Optional<FolioEntity> folio =

folioRepository.findByFolioIdLocked(event.getObjectId());

        if (folio.isEmpty()) {
            folioEntity = mapper.createFolioEntity(event);
        } else {
            folioEntity = mapper.updateFolioEntity(event,
folio.get());
        }
    });
}
```

```
        }  
        return folioRepository.save(folioEntity);  
    })  
    .flatMap(folioRepositoryService::addTransientFields);  
}
```

## GENERIC DB SHARED LOCK TABLE

```
public <T> Mono<T> acquireLock(Long sourceId, String reportableId,
    LockType lockType, Supplier<Mono<T>> callbackOnLockAcquisition) {
    return Mono.defer(
        () ->
            sharedLockRepository
                .findByLockTypeAndSourceIdLocked(lockType,
sourceId)
                .or(() -> Optional.of(
sharedLockRepository.save(SharedLockEntity.builder()
    .lockType(lockType).sourceId(sourceId).failureCount(0)
        .expiryCount(0).build()))))
                .flatMap(
                    lock -> {
                        if (lock.getAcquiredLockTime() != null)
return Optional.empty();
                        lock.setAcquiredLockTime(Instant.now());
                        val updatedLock =
sharedLockRepository.save(lock);
                        return Optional.of(updatedLock);
                    })
                .onErrorResume( // OpLock exeception for existing rows,
integrity violation for conflicting row inserts
                    t -> t instanceof ObjectOptimisticLockingFailureException
|| t instanceof DataIntegrityViolationException,
                    e -> Mono.empty())
    )
}
```



```
        })
        .flatMap(lock -> callbackOnLockAcquisition.get()) // We have
        acquired the lock, call the callback
        .flatMap(
            resultItem ->
                tenantContext.fromSupplier(
                    () -> {
                        log.info("Remove lock for {}, sourceId: {}",
reportableId, sourceId);
                        sharedLockRepository.unsetLock(lockType,
sourceId);
                        return resultItem;
                    })
                ));
    }
```

## **5 MINS != EXPERT :)**

- To lock or not to lock
  - Complexity trade-off
    - DB locks are simple until it's use grows
  - Idempotent updates
  - Event normalisation
- Audit trail / logging

## QUESTIONS

- <https://github.com/andrewrembrandt/distributed-locks-talk>
- <https://blog.mimacom.com/tag/concurrency/>